

```
In [1]: import numpy as np
```

```
In [2]: import pandas as pd
```

```
In [3]: from matplotlib import pyplot as plt
```

```
In [4]: import seaborn as sb
```

```
In [5]: from sklearn.model_selection import train_test_split
```

```
In [6]: from sklearn.preprocessing import MinMaxScaler
```

```
In [8]: from sklearn import metrics
```

```
In [9]: from sklearn.svm import SVC
```

```
In [10]: from xgboost import XGBClassifier
```

```
In [11]: from sklearn.linear_model import LogisticRegression
```

```
In [12]: import warnings  
warnings.filterwarnings('ignore')
```

```
In [13]: df=pd.read_csv('WineQT.csv')
```

```
In [15]: print(df.head())
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality	Id
0	9.4	5	0
1	9.8	5	1
2	9.8	5	2
3	9.8	6	3
4	9.4	5	4

```
In [16]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1143 entries, 0 to 1142
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1143 non-null   float64
1   volatile acidity       1143 non-null   float64
2   citric acid            1143 non-null   float64
3   residual sugar         1143 non-null   float64
4   chlorides              1143 non-null   float64
5   free sulfur dioxide    1143 non-null   float64
6   total sulfur dioxide   1143 non-null   float64
7   density               1143 non-null   float64
8   pH                    1143 non-null   float64
9   sulphates             1143 non-null   float64
10  alcohol               1143 non-null   float64
11  quality               1143 non-null   int64
12  Id                   1143 non-null   int64
dtypes: float64(11), int64(2)
memory usage: 116.2 KB

```

In [18]: `df.shape`

Out[18]: (1143, 13)

In [19]: `df.describe()`

Out[19]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000
mean	8.311111	0.531339	0.268364	2.532152	0.086933	15.615486	45.914698
std	1.747595	0.179633	0.196686	1.355917	0.047267	10.250486	32.782130
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
25%	7.100000	0.392500	0.090000	1.900000	0.070000	7.000000	21.000000
50%	7.900000	0.520000	0.250000	2.200000	0.079000	13.000000	37.000000
75%	9.100000	0.640000	0.420000	2.600000	0.090000	21.000000	61.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	68.000000	289.000000

In [20]: `df.isna().any`

```

Out[20]: <bound method NDFrame._add_numeric_operations.<locals>.any of      fixed acidity
volatile acidity citric acid residual sugar chlorides \
0                False                False                False                False                False
1                False                False                False                False                False
2                False                False                False                False                False
3                False                False                False                False                False
4                False                False                False                False                False
...                ...                ...                ...                ...                ...
1138             False                False                False                False                False
1139             False                False                False                False                False
1140             False                False                False                False                False
1141             False                False                False                False                False
1142             False                False                False                False                False

      free sulfur dioxide total sulfur dioxide density    pH sulphates \
0                False                False    False    False    False
1                False                False    False    False    False
2                False                False    False    False    False
3                False                False    False    False    False
4                False                False    False    False    False
...                ...                ...    ...    ...    ...
1138             False                False    False    False    False
1139             False                False    False    False    False
1140             False                False    False    False    False
1141             False                False    False    False    False
1142             False                False    False    False    False

      alcohol  quality    Id
0        False    False    False
1        False    False    False
2        False    False    False
3        False    False    False
4        False    False    False
...        ...    ...    ...
1138     False    False    False
1139     False    False    False
1140     False    False    False
1141     False    False    False
1142     False    False    False

[1143 rows x 13 columns]>

```

```
In [21]: df.isnull().sum()
```

```

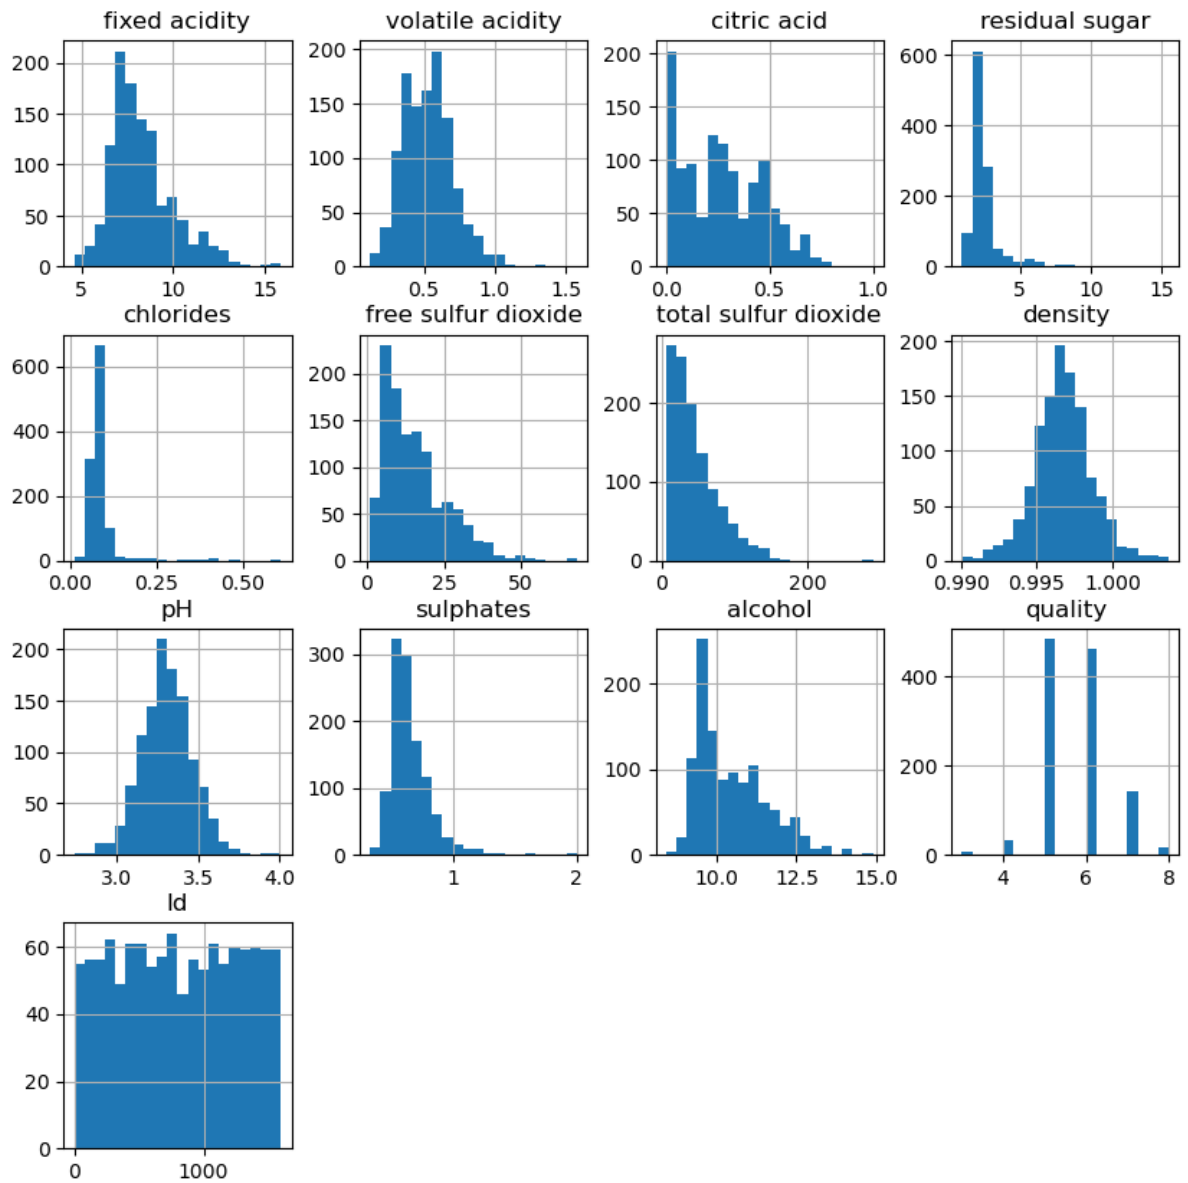
Out[21]: fixed acidity      0
volatile acidity    0
citric acid         0
residual sugar      0
chlorides           0
free sulfur dioxide  0
total sulfur dioxide 0
density             0
pH                 0
sulphates           0
alcohol             0
quality             0
Id                 0
dtype: int64

```


```
In [22]: df.columns
```

```
Out[22]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
        'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
        'pH', 'sulphates', 'alcohol', 'quality', 'Id'],  
        dtype='object')
```


```
In [23]: df.hist(bins=20,figsize=(10,10))  
plt.show()
```



```
In [24]: plt.bar(df['quality'],df['alcohol'])  
plt.xlabel('quality')  
plt.ylabel('alcohol')  
plt.show()
```



```
In [26]: plt.figure(figsize=(12,12))
sb.heatmap(df.corr()>0.6,annot=True,cbar=False)
plt.show()
```



In []: *#From the above heat map we can conclude that the 'total sulphur dioxide' and 'free*

```
In [28]: df=df.drop('total sulfur dioxide',axis=1)
```

```
In [29]: #model development
```

```
In [30]: df['best quality'] = [1 if x > 5 else 0 for x in df.quality]
```

```
In [31]: df.replace({'white': 1, 'red': 0}, inplace=True)
```

```
In [32]: features = df.drop(['quality', 'best quality'], axis=1)
target = df['best quality']
```

```
xtrain, xtest, ytrain, ytest = train_test_split(
    features, target, test_size=0.2, random_state=40)
```

```
xtrain.shape, xtest.shape
```

Out[32]: ((914, 11), (229, 11))

```
In [33]: norm = MinMaxScaler()
xtrain = norm.fit_transform(xtrain)
xtest = norm.transform(xtest)
```

```
In [34]: models = [LogisticRegression(), XGBClassifier(), SVC(kernel='rbf')]

for i in range(3):
    models[i].fit(xtrain, ytrain)

    print(f'{models[i]} : ')
    print('Training Accuracy : ', metrics.roc_auc_score(ytrain, models[i].predict(>
    print('Validation Accuracy : ', metrics.roc_auc_score(
        ytest, models[i].predict(xtest)))
    print()


LogisticRegression() :
Training Accuracy : 0.7546950559364851
Validation Accuracy : 0.7255154639175256

XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=None, n_jobs=None,
               num_parallel_tree=None, random_state=None, ...) :
Training Accuracy : 1.0
Validation Accuracy : 0.8022102467978757

SVC() :
Training Accuracy : 0.7648213641284736
Validation Accuracy : 0.7358247422680412
```

```
In [ ]: #Model Evaluation
        From the above accuracies we can say that Logistic Regression and SVC() classifier
```

```
In [35]: metrics.plot_confusion_matrix(models[1], xtest, ytest)
         plt.show()
```



```
In [36]: print(metrics.classification_report(ytest,
                                             models[1].predict(xtest)))
```

	precision	recall	f1-score	support
0	0.75	0.79	0.77	97
1	0.84	0.81	0.83	132
accuracy			0.80	229
macro avg	0.80	0.80	0.80	229
weighted avg	0.81	0.80	0.80	229

```
In [ ]:
```